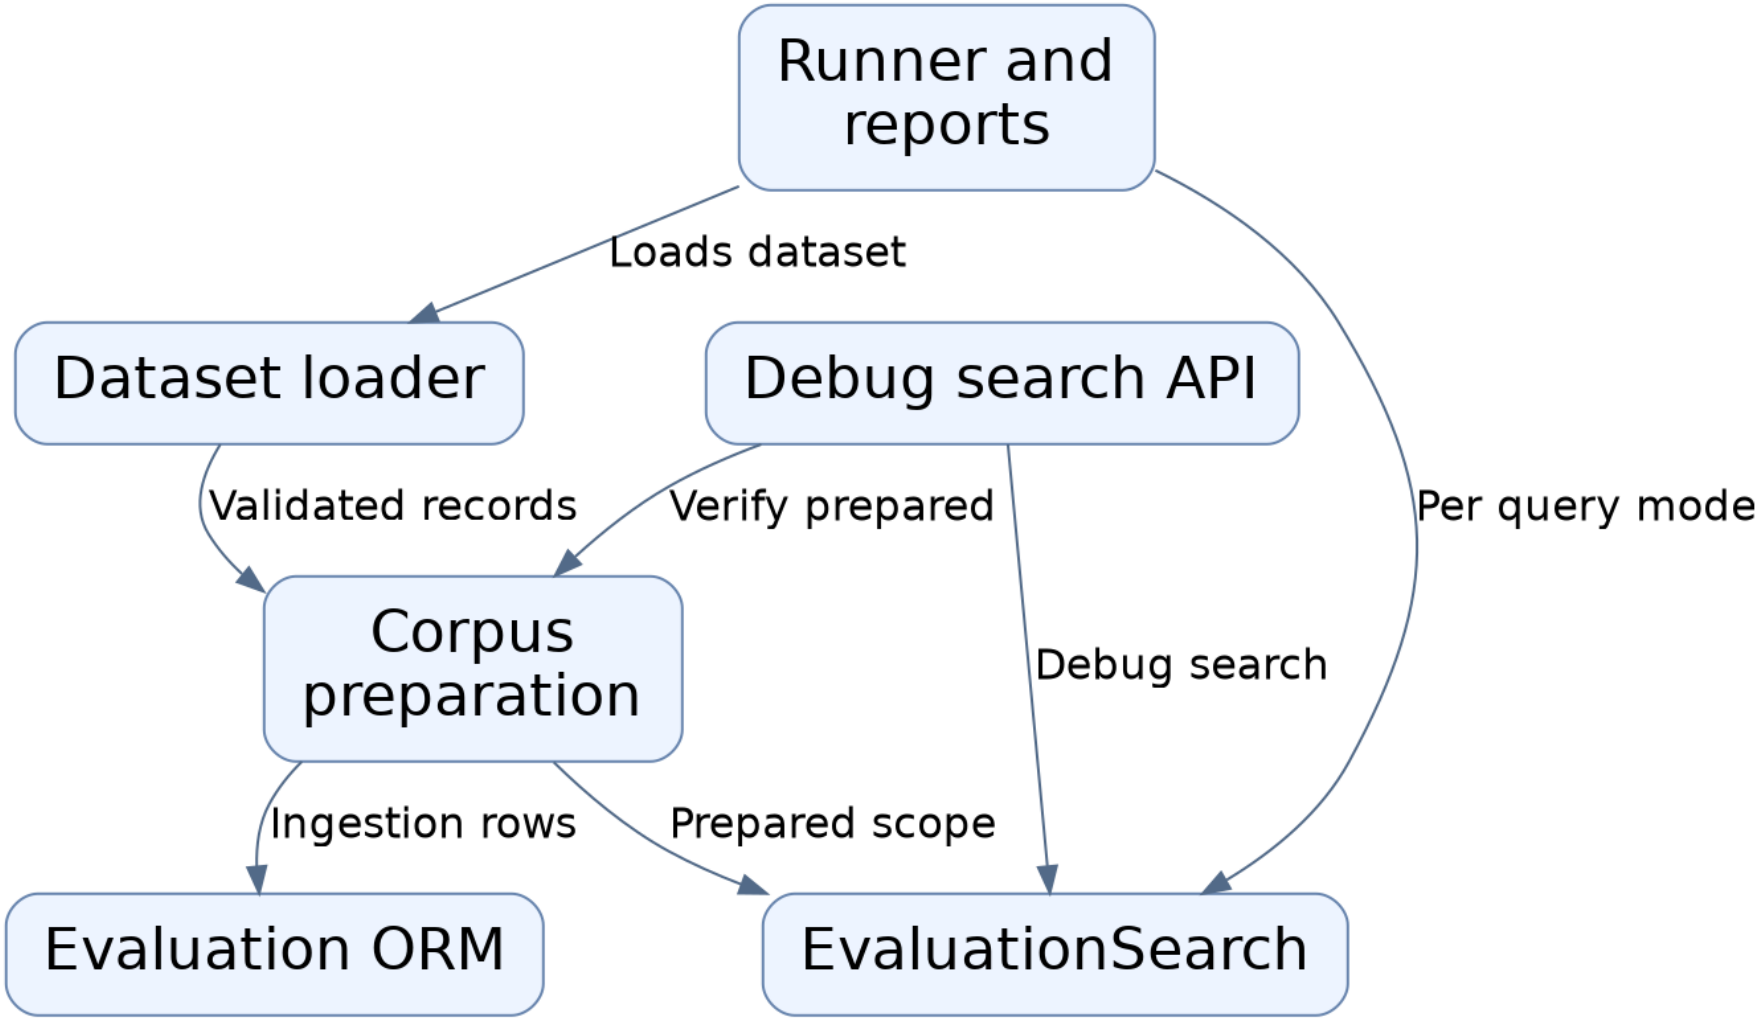
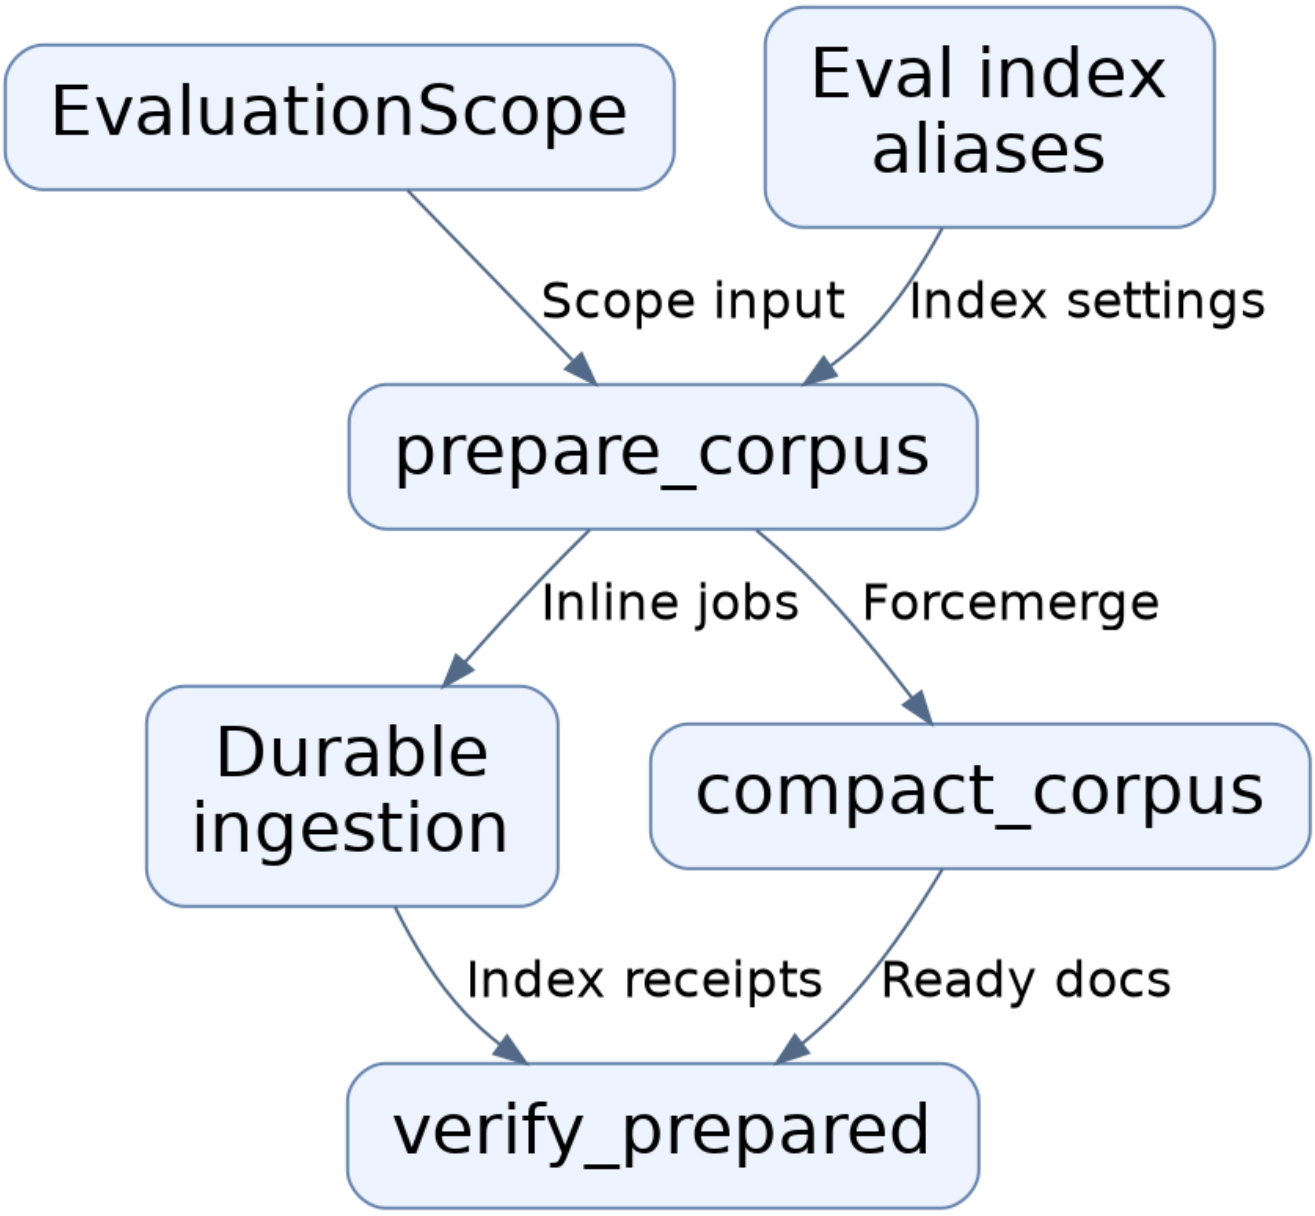
Architecture

Evaluation composes dataset loading, corpus preparation, mode-dispatch search, scoring/reporting, its HTTP and CLI entry points, and persistence; boundaries matter because corpus state and retrieval faults map to distinct errors.



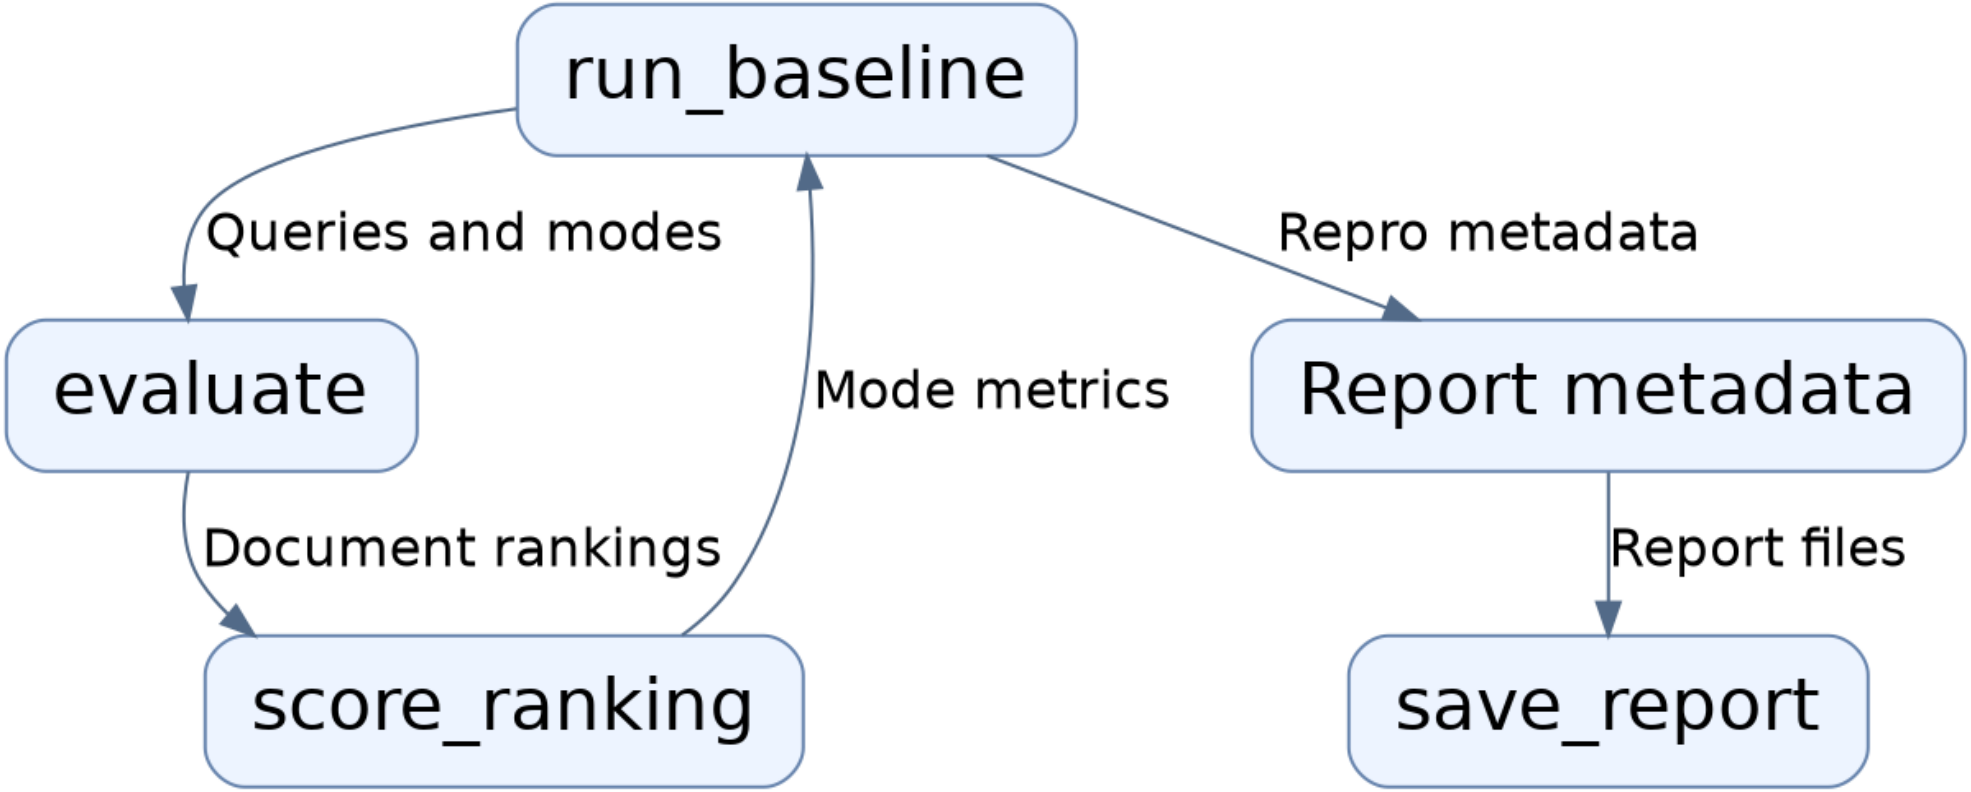
Corpus preparation flow

prepare_corpus derives a deterministic scope and index aliases, reuses the durable ingestion pipeline inline, treats persisted documents as immutable, then force-merges and verifies against job receipts.



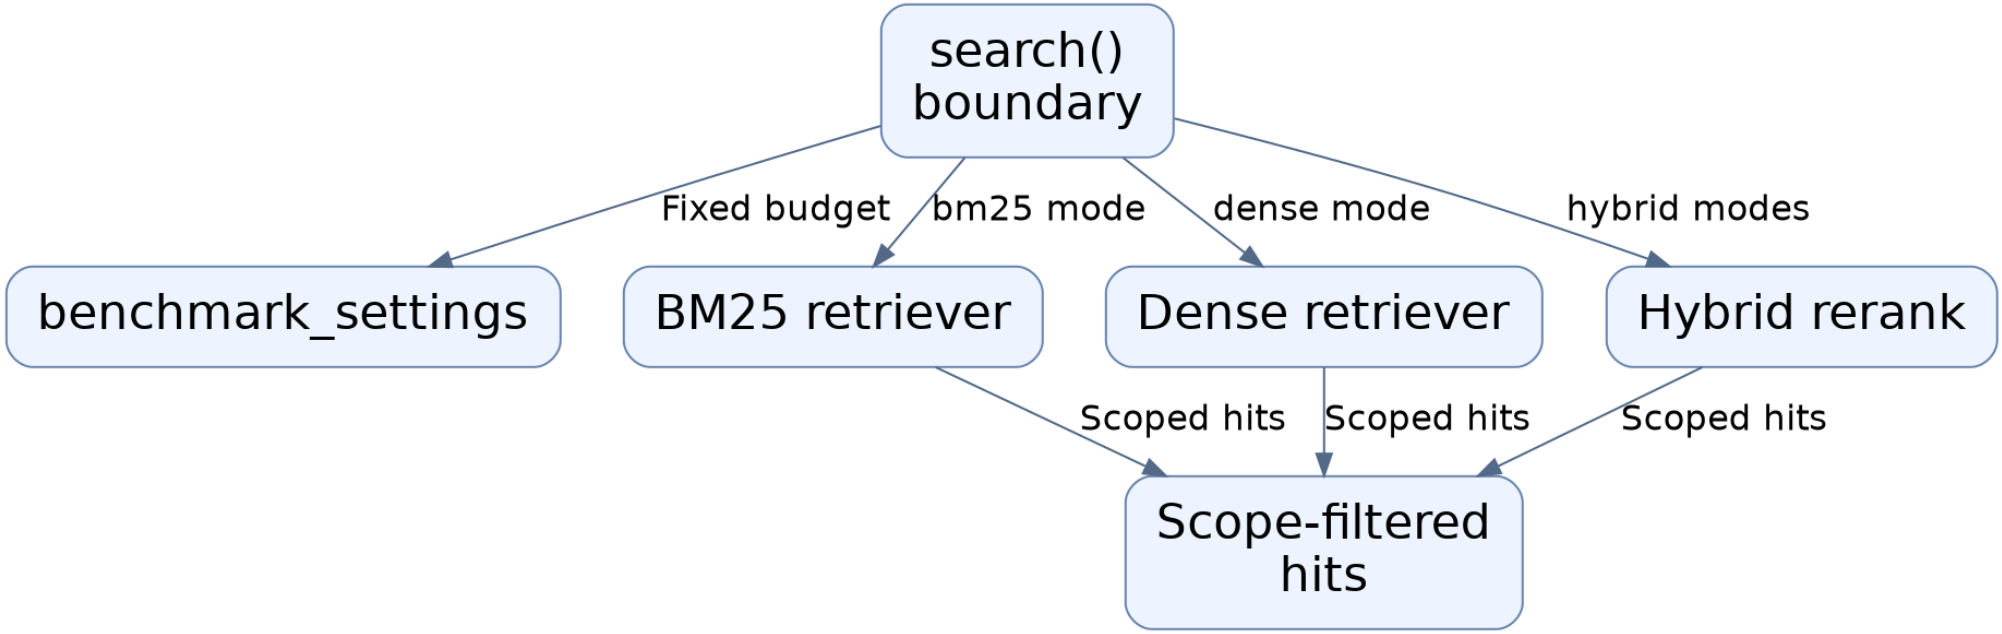
Evaluation run flow

run_baseline selects and prepares a dataset, calls evaluate across modes, scores each ranking against judgments, attaches credential-free reproduction metadata, and writes JSON plus markdown reports.



Search mode dispatch

EvaluationSearch validates mode and filters, fixes a benchmark budget, lazy-builds the mode-specific retriever, then asserts every result stays inside the resolved scope before deduplicating documents.



SciFact dataset build

The SciFact CLI verifies a pinned archive download, converts BEIR files into the shared four-file contract, and can build a deterministic 30-query/500-doc subset that load_dataset re-validates.

